COL106: Data Structures, I Semester 2018-19
Assignment 3-4
A small search engine

September 24, 2018

In this assignment we will build the basic data structure underlying search engines: an *inverted index*. We will use this inverted index to answer some simple search queries.

An inverted index for a set of webpages

Suppose we are given a set of webpages W. For our purposes, each webpage $w \in W$ will be considered to be a sequence of words $w_1, w_2, \ldots w_k$. Another way of representing the webpage could be to maintain a list of words along with the position(s) of the words in the webpage. For example consider a web page with the following text:

Data structures is the study of structures for storing data.

This can be represented as

$$\{(\text{data} : 1, 10), (\text{structures} : 2, 7), (\text{study} : 5), (\text{storing} : 9)\}.$$

Note that the small connector words like "is", "the", "of", "for" have not been stored. Words like this are referred to as *stop words* and are generally removed since they are very frequent and normally contain no information about the content of the webpage.

This representation of the webpage is similar to the index we see at the back of many books which tell us the page numbers where certain important terms used in the book may be found. In fact, we can refer to this as an

index for the webpage. In mathematical notation we would say that given a webpage $w = w_1, w_2, \ldots, w_k$, the index of w is

$$\{(u : i_1(u), \ldots i_\ell(u)) : w_{i_j(u)} = u, 1 \leq j \leq \ell\}.$$

An index is used to find the location of a particular string (word) in a specific document or webpage, but when we move to a *collection* of webpages, we need to first figure out which of the web pages contain the string. For this we store an *inverted index*. Let us try to define an inverted index formally.

Let us suppose we are given a collection $\mathcal{C}$ of webpages. For each page $p \in C$, let us denote by $W(p)$ the set of all words (excluding stop words) that occur in p. Note that

$$W(\mathcal{C}) = \bigcup_{p \in \mathcal{C}} W(p),$$

is the set of all words in our collection.

An inverted index for $\mathcal{C}$ will contain an entry for each word $w \in W(\mathcal{C})$. This entry will contain tuples of the form (p, k) to indicate that w occurs in the kth position of page $p \in \mathcal{C}$. Using the notation that $p[k]$ denotes the kth word of page p, we can say that the inverted index of $\mathcal{C}$ is defined as

$$\text{Inv}(\mathcal{C}) = \{(w : \{(p, k) : p \in \mathcal{C}, p[k] = w\}) : w \in W(\mathcal{C})\}.$$

For example, consider the following (small) collection of documents.

1: Data structures is the study of structures for storing data.

2: Structural engineers collect data about structures

The inverted index for this would be

```
{(data : {(1,1), (1,10), (2,4)}),
(structures: {(1,2), (1,7), (2,6)}),
(study : {(1,5)}),
(storing : {(1,9)}),
(structural : {(2,1)}),
(engineers : {(2,2)}),
(collect : {(2,3)}) }
```

The web search problem

The *web search problem* is defined as follows:

> Given a collection of webpages $\mathcal{C}$ and a sequence of words $q_1 \ldots q_k$, find the "most relevant" set of pages $p_1, p_2, \ldots p_\ell$ that contain as many of $q_1 \ldots q_k$ as possible and return them in the order of decreasing "relevance."

The question of how to measure the relevance of a webpage to a particular query is an involved question with no easy answers. However, for the purpose of this assignment we will work with a simple scoring function.

A scoring function for search term relevance

One of the simplest scoring function is term frequency-inverse document frequency. It is used to measure how important a word w is to a webpage p. It is a product of two factors i.e. term frequency and inverse document frequency. Given a word w and a webpage p, the relevance score is defined as

TF-IDF

$$\text{relevance}_w(p) = tf_w(p) * idf_w(p)$$

Term Frequency: It is the total number of occurrence of a word w in a webpage p, denoted by $f_w(p)$. It is normalized by the total number of words in webpage p, denoted by $|W(p)|$. It is defined as

$$\text{tf}_w(p) = \frac{f_w(p)}{|W(p)|}$$

Inverse document frequency: It is the logarithm of the total number of webpages, denoted by N divide by the logarithm of the number of webpages $n_w(p)$ that contain the word w. It is defined as

$$\text{idf}_w(p) = \log \frac{N}{n_w(p)}$$

So, if we are given a search query that has a single term, say w, to return the webpages in order of relevance we have to first extract the entry corresponding to w from $\text{Inv}(\mathcal{C})$ and then calculate the relevance of each page and return the pages in decreasing order of relevance.

Compound searches

In this assignment we will answer three kinds of search queries: AND queries, OR queries and phrase queries. We now describe these three along with their scoring methodology.

- **OR** queries: Given a search query $q_1 \ldots q_k$, any page that contains *any* of the words q_1 to q_k is a valid answer. The relevance score of a page p is computed as

$$\text{relevance}_{q_1 \ldots q_k}(p) = \sum_{i=1}^{k} \text{relevance}_{q_i}(p),$$

  and pages are returned in decreasing order of relevance. Note that if some q_i does *not* occur in page p the $\text{relevance}_{q_i}(p) = 0$.

- **AND** queries: Given a search query $q_1 \ldots q_k$, any page that contains *all* of the words q_1 to q_k is a valid answer. The relevance score of a page p is computed as

$$\text{relevance}_{q_1 \ldots q_k}(p) = \sum_{i=1}^{k} \text{relevance}_{q_i}(p),$$

  and pages are returned in decreasing order of relevance.

- **Phrase** queries : Given a search query $q_1 \ldots q_k$, any page that contains q_1 in position ℓ, q_2 in position $\ell+1$ and so on till q_k in position $\ell+k-1$ is said to contain the phrase $q_1 \ldots q_k$ at the position ℓ. Suppose in a webpage p having $|W(p)|$ words, the phrase $q_1 \ldots q_k$ occurs m times then the relevance score of page p for this phrase is

$$\text{relevance}_{q_1 \ldots q_k}(p) = tf_{q_1 \ldots q_k}(p) * idf_{q_1 \ldots q_k}(p)$$

$$= \frac{m}{|W(p)| - (k-1) * m} * \log \frac{N}{n_w(p)}$$

The programming assignment

The implementation of the search engine has been divided into two assignments. You have been provided with separate folders for each assignment

having their own webpages folder and actions.txt. Use the checker files from earlier assignments to read the action files if you need to.

In the first part of the assignment, Assignment 3, we will build the basic backbone of the search engine and answer single word search queries. The answers to the actions.txt for this parts are given in answers.txt. *There is no auto-grading. Compare your output to the output in answers.txt. You can also make your own data files and action files and run your code on that.*

In the second part of the assignment, Assignment 4, compound queries will be implemented. *Note that for this there is no answers.txt although we have provided an actions file. You have to verify the correctness of your algorithm/program by yourself.* Both the assignments have different deadlines.

Assignment 3 [100 marks]
Deadline: 11:55PM, 13 October 2018

- Write a Java class `MySet` using Java generic's
  (*https://docs.oracle.com/javase/tutorial/java/generics/types.html*).
  The class should be represented as *MySet<X>* where X is the datatype of the set. MySet should implement the following methods:

  - `void addElement(X element)`: Add `element` to the set.
  - `MySet<X> union(MySet<X> otherSet)`: Return MySet which represents a union of the current set and the `otherSet`.
  - `MySet<X> intersection(MySet<X> otherSet)`: Return MySet which represents an intersection of the current set and the `otherSet`.

- Write a Java class `MyLinkedList` using Java generic's. It should contain the standard methods of a linked list.

- Write a Java class `Position` that represents a tuple *<page p, word position i>*.

  - `Position(PageEntry p, int wordIndex)` Constructor method.
  - `PageEntry getPageEntry()` Return p
  - `int getWordIndex()` Return wordIndex

- Write a Java class `WordEntry`. For a string *str*, this class stores the list of word indice's where str is present in the document(s).

- **WordEntry(String word)**: Constructor method. The argument is the word for which we are creating the word entry.
- **void addPosition(Position position)**: Add a position entry for str.
- **void addPositions(MyLinkedList<Position> positions)**: Add multiple position entries for str.
- **MyLinkedList<Position> getAllPositionsForThisWord()**: Return a linked list of all position entries for str.
- **float getTermFrequency(String word)**: Return the term frequency of the word in a webpage.

- Write a Java class **PageIndex** which stores one word-entry for each *unique* word in the document.

  - **void addPositionForWord(String str, Position p)**: Add position p to the word-entry of **str**. If a word entry for **str** is already present in the page index, then add p to the word entry. Otherwise, create a new word-entry for **str** with just one position entry p.
  - **LinkedList<WordEntry> getWordEntries()**: Return a list of all word entries stored in the page index.

- Write a Java class **PageEntry** to store the the information related to a webpage. It should contain following methods:

  - **PageEntry(String pageName)**: Constructor method. The argument is the name of the document. Read this file, and create the page index.
  - **PageIndex getPageIndex()**: This method returns the page index of this web-page.

- Write a Java class **MyHashTable** that implements the hashtable used by the InvertedPageIndex. It maps a word to its word-entry.

  - **private int getHashIndex(String str)**: Create a hash function which maps a string to the index of its word-entry in the hashtable. The implementation of hashtable should support chaining.

6

– void addPositionsForWord(WordEntry w): This adds an entry to the hashtable: $stringName(w) -> positionList(w)$. If no word-entry exists, then create a new word entry. However, if a word-entry exists, then merge w with the existing word-entry.

- Write a Java class InvertedPageIndex which contains the following methods:

  – void addPage(PageEntry p): Add a new page entry p to the inverted page index.

  – MySet<PageEntry> getPagesWhichContainWord(String str): Return a set of page-entries of webpages which contain the word str.

- Write a Java class SearchEngine. This is the class that we will use as an interface to the search engine. It should contain following methods:

  – SearchEngine(): This is the constructor method. It should create an empty InvertedPageIndex.

  – void performAction(String actionMessage): This the main stub method that you have to implement. It takes an action as a string. The list of actions, and their format will be described later.

Actions:

- addPage x Add webpage x to the search engine database. The contents of the webpage are stored in a file named x in the webpages folder.

- queryFindPagesWhichContainWord x Print the name of the webpages which contain the word x. The list of webpage names should be comma separated. If the word is not found in any webpage, then print "No webpage contains word x"

- queryFindPositionsOfWordInAPage x y Print the word indice's where the word x is found in the document y. The word indices should be separated by a comma. If the word x is not found in webpage y, then print "Webpage y does not contain word x". If the webpage is not added in database, then print "No webpage y found".

Points to note:

- Convert each word to lowercase.

- Do not store the connector words in the search engine. However, consider them when you calculate the word indice's. Here is a list of connector words: { *a, an, the, they, these, this, for, is, are, was, of, or, and, does, will, whose* }.

- Replace the punctuation marks with a *space*. Here is a list of punctuations: { } [ ] < > = ( ) . , ; ' " ? # ! - :

- *Plural and singular form*: Assume that these words are same: (stack, stacks), (structure, structures), (application, applications).

- We have given you a set of *connector words*, *punctuation marks*, and *singular-plural entries*. Consider these list of words as exhaustive. You do not need to make any more exceptions in your code.

- actions.txt contains some simple search queries.

- answers.txt contains the answers to actions.txt for you to check your code. *There is no autograding for this assignment.*

Assignment 4 [100 marks]
Deadline: 11:55PM, 27 October 2018

In this assignment, we will complete the implementation of our search engine.

- Implement the following method in the `InvertedPageIndex` class.

    - `MySet<PageEntry>getPagesWhichContainPhrase(String str[])`: Return a set of page-entries for webpages which contain a sequence of *non-connector* words (`str[0] str[1] ... str[str.len-1]`). Assume a webpage which contains the following text: *"Data structures is the study of structures for storing data."* This webpage contains the phrases: *"Data structures"*, *"Data structures study"*, and *"Data structures study structures"*.

- Implement the collection of position entries for the `WordEntry` class as an AVL-tree which is a height-balanced binary search tree. The data item for each node in the the tree is a *position entry*, and the ordering of the nodes in the tree is on the basis of the word index in the position entry. For phrase queries, you must use the AVL tree to find the next word.

- Implement the following method in the `PageEntry` class.

  - `float getRelevanceOfPage(String str[ ], boolean doThese-WordsRepresentAPhrase)`: Return the relevance of the webpage for a group of words represented by the array `str[ ]`. If the flag `doTheseWordsRepresentAPhrase` is true, it means that the words represent a phrase; otherwise the words are part of a complex query (AND/OR).

- Write a Java class `SearchResult` which represents a tuple<*page p, relevance r*>. `SearchResult` implements the Java interface `Comparable` (http://docs.oracle.com/javase/7/docs/api/java/lang/Comparable.html).

  - `public SearchResult(PageEntry p, float r)`: Constructor method.
  - `public PageEntry getPageEntry()`: Return p.
  - `public float getRelevance()`: Return r.
  - `public int compareTo(SearchResult otherObject)`: Gives the ordering between the current object and the `otherObject`.

- Write a Java class `MySort` to sort a list of `Comparable` objects. It must contain one method:

  - `ArrayList<Sortable> sortThisList(MySet<Sortable> listOf-SortableEntries)`: Given a set of `Sortable` objects, this method returns a sorted list of objects. The list is represented as Java's `ArrayList` where the following relation holds: if a<b, *sortedlist.get(a).getNumber() >= sortedlist.get(b).getNumber()*. You can implement any sorting algorithm that you want. Your SearchEngine class should use the MySort class to sort the set of pages on the basis of the relevance criteria.

Actions:

- `queryFindPagesWhichContainAllWords x1 x2 .. xn`: Print the name of the webpages which contain all the words `x1, x2, .. xn`. The words are separated by a space.

- `queryFindPagesWhichContainAnyOfTheseWords x1 x2 .. xn`: Print the name of the webpages which contain at least one word from this set { `x1, x2, .. xn` }.

- `queryFindPagesWhichContainPhrase x1 x2 .. xn`: Print the name of the webpages which contain the phrase `x1 x2 .. xn`.

Points to note:

- The search engine should print the name of all the webpages which satisfy the given criteria. The list of webpages should be displayed in a sorted order.

- Assume that the input to the search engine does not contain any *punctuation mark* or a *connector word*.

- actions.txt contains some compound search queries.

- *No answers.txt corresponding to actions.txt is provided. There will be no autograding either.*

- Actions available for assignment 3 should also be included in assignment 4.